

Class 06

Shared Folders on Kali Linux

SSDs

Volumes

Partitions

Shared Folders on Kali Linux

**Sharing Folders Between RADISH
(Windows 10) and your Kali Linux VM**

Background

From Last Week

Last lecture, we set up Kali Linux on your RADISH desktop

Be sure you have completed this task by Friday, 30 Sept

During that exercise, we set up shared folders between the RADISH (Windows 10) VM and the Kali Linux VM

We shared three folders (available on RADISH):

E: drive

R:\share

Your documents directory

The shared folders SHOULD be available at the following directory location on Kali Linux:

/mnt/hgfs

Sharing Issue Resolution

However, you may find that the **hgfs** folder doesn't exist in your **/mnt** directory

The following slides will walk you through the resolution

The Fix

Page 1 of 4

1. Power up the Kali Linux VM and log in
login: kali
password: kali
2. On Workstation Pro, using the **VM > Settings** menu option to open up the Virtual Machine Settings window
3. Click on the **Options** in the Virtual Machine Settings window
4. Click on **Shared Folders** on the left pane
5. Click the **Disabled** radio button the upper right of the window
6. Click **OK** on the lower right

The Fix

Page 2 of 4

7. Open up a Terminal window in your Kali Linux VM
8. Type in the following two commands:
`sudo apt-get remove open-vm-tools`
`sudo apt-get purge open-vm-tools`
9. Reboot Kali Linux
10. Once Kali Linux has finished rebooting, log in
11. Select the following Workstation Pro menu item:
VM > Install VMWare Tools
12. You will then see a message bar at the bottom of your Kali Linux windows telling you to install VMWare Tools via a CD

The Fix

Page 3 of 4

13. Double-click on the **File System** icon on your Kali desktop to open the **File Manager**
14. Click on **Computer** on the left panel of the **File Manager**
15. Verify that you see “VMWare Virtual IDE CDRom Drive” displayed on the right
16. Open up a Terminal window in your Kali Linux VM
17. Go to the temp directory by using the following command:
`cd /tmp`
18. Unzip the VMWare tools package into your tmp directory using the following command:
`tar xzpf /medi/kali/VMkware\ Tools/VMwareTools-10.3.23-16594550.tar.gz`
19. Navigate to the tools distribution directory with the following command:
`cd vmware-tools-distrib`
20. Run the VMware Tools installation script with the following command:
`sudo ./vmware-install.pl`
21. Follow the prompts to accept the default values.
22. Ensure the script completes execution

The Fix

Page 4 of 4

23. Shut down Kali Linux
24. Exit Workstation Pro
25. Bring up Workstation Pro
26. Power up Kali Linux
27. Log in to Kali Linux
28. On Workstation Pro, using the **VM > Settings** menu option to open up the Virtual Machine Settings window
29. Click on the **Options** in the Virtual Machine Settings window
30. Click on **Shared Folders** on the left pane
31. Click the **Always enabled** radio button the upper right of the window
32. Click **OK** on the lower right
33. In the Kali VM, double-click on the File System icon on the desktop
34. Verify that you have the **hgfs** folder in your **/mnt** directory

Flash Memory & SSD File Systems

Overview

Flash/SSD Memory

USB Flash/SSD Drives

Flash/SSD File Systems

Relation to conventional file systems

Flash/SSD Memory

Overview

History

Basic Information

NAND Flash & SSD Memory Organization

Some History

Invented in early 1980s by Toshiba

Called “flash” because erasing the memory reminded scientists of the flash of a camera

Intel offered the first NOR (Not OR) flash memory products in late 1980s

Keeps its memory state without needing power

Read/write random access; state can be changed electronically

Code that is stored can be executed directly from NOR flash

Targeted at replacing ROMs (not changeable) and EPROMs (erased using UV light or X-ray)

Useful for electronic configuration info and read-only code such as for BIOSs and TV set-top boxes

Some History

NAND flash memory announced by Toshiba in mid 1980s

Much higher bit density and write speeds than NOR flash

Could be accessed much like magnetic disk drives

*Was thought of as being an eventual replacement for mass storage
such as magnetic disk drives*

Drawback: Reads & writes had to be done in blocks

First removable NAND flash memory in mid 1990s

Looked like a large SD card

First NAND USB drive in 2000

Until the advent of smart phones, NAND flash memory was used
mostly for USB flash drives and camera SD cards

Some History

2000: USB 2.0 standard specified 1.5, 12 or 480 Mbps

But NAND flash could do only 100 Mbps reads


A lot less for writes

2008: USB 3.0 standard specified up to 5 Gbps

Even today's faster NAND flash memory can't come close

Some Basic Flash/SSD Info

Electronic Evolution*



Vacuum tube triode

Thermion amplifier

Invented: 1906 by Lee de Forest

Later on IIT's Armour

College physics faculty

Originally used

Phone systems


Radio & TV

Very early computers

*e.g., AVIDAC**

Uses today

Power RF transmitters



Junction transistor

Created: 1947

*W.Brattan, J.Bardeen,
W.Schockley at Bell Labs*

Nobel laureates

Used


Phone systems

Radio & TV

Computers

Uses today

*Electronic power
devices*



Field-Effect Transistor

MOSFET

Created: 1959

*D.Kahng, M.Atalla at
Bell Labs*

Uses today

*Almost all electronic
devices and systems*

Flash Bit Cell

float gate (FG) completely insulated by metal oxide


Electrons put on *FG* can stay there for years

The electron charge on the *FG* affects the field imposed by the CG

Reading the bit

If there is a negative charge on the FG, less current will flow in the bit line for a given CG-to-source voltage

And the converse if there is no FG charge



Putting Charge on the FG


The normal charge on the FG is 0
To put a negative charge on the FG:

Apply a high $+V_{GS}$ voltage to the CG with the bit line also at a high voltage

Because there is no charge on the FG, many electrons will flow from the source to the bit line


Some electrons will have sufficient energy to tunnel through the insulating layer to the FG

Called **hot injection**



Removing Charge from the FG

The source is disconnected
A large + voltage is applied to the bit line
A large negative voltage relative to the bit line is applied to the CG
This causes electrons to be attracted to the N-doped silicon connected to the bit line
If the voltage is high enough, electrons will leave the FG by tunneling through the insulation



Bit Wear

Reading the bit that is stored on the float gate is done by simply sensing the bit line current

No charge is added or removed from the float gate

Reading can be done an almost infinite number of times with no degradation

But changing the charge on the float gate involves hot injection

This causes wear on the float gate insulation

After a number of program/erase (P/E) cycles, the metal oxide surrounding the FG will become degraded

Then the FG is unable to reliably hold a charge

NOR Flash

Called NOR because the transistors are arranged as a NOR gate


If any word line is high (1), it's flash cell conducts

This causes the bit line voltage to be close to ground, i.e., low (0)

A high output (1) is produced only if all inputs are low (0)

Here is a NOR truth table for two inputs (words)

Input		Output
A	B	A NOR B
0	0	1
0	1	0
1	0	0
1	1	0



NAND Flash

Here the FETs are arranged as a NAND gate


AND: $WL0 \& WL1 \& \dots WL7 = 1$

NAND: $WL0 \& WL1 \& \dots WL7 = 0$

All word lines, the bit select and the ground select must be high in order for the bit line to be low

Here is the NAND truth table for two inputs (words)

Input		Output
A	B	A NAND B
0	0	1
0	1	1
1	0	1
1	1	0



NAND Flash

NAND flash memory density can be made much greater than NOR flash. How?

Reduced the number of wires relative to NOR flash

Only the word lines and a single ground and bit line are needed to access a bit

Multi-level cells

NAND flash is arranged so as to not access individual bytes

Taking advantage of NAND structure

NAND flash is arranged so that

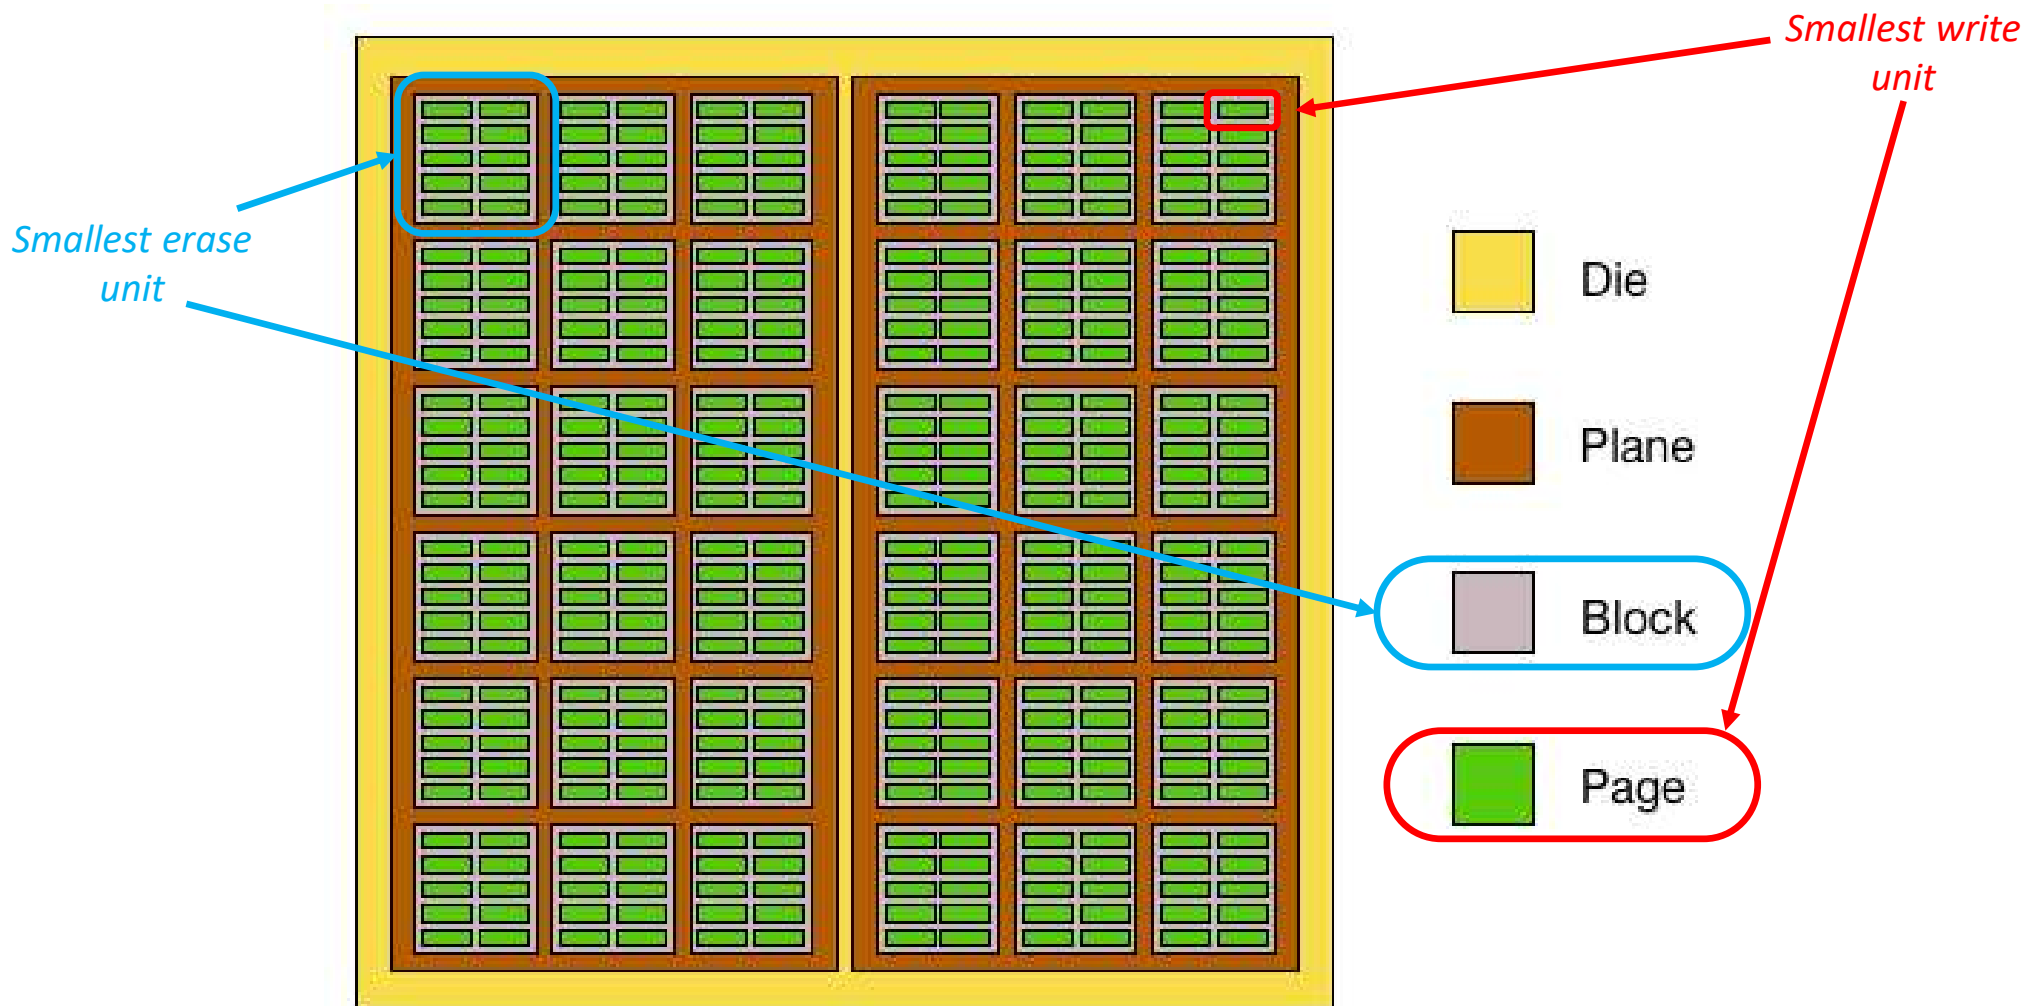
*It can be written as a group of bytes (e.g., a **page**) at a time*

*But it can only be erased a **block** at a time*

Blocks >> pages

NAND Flash

Blocks and Pages



Multi-level NAND Flash Cells

There are currently 4 types of NAND flash cells

SLC (single-level cell): 1 bit per cell

MLC (multi-level cell): 2 bits per cell

TLC (tri-level cell): 3 bits per cell

QLC (quad-level cell): 4 bits per cell

How can we get more than 1 bit per cell?

Multi-level NAND Flash Cells

How?

We can put different charges on the FG

SLC: 1 bit per cell \Rightarrow 2 state of charge on the FG

MLC: 2 bits per cell \Rightarrow 4 states of charge on the FG

TLC: 3 bits per cell \Rightarrow 8 states of charge on the FG

QLC: 4 bits per cell \Rightarrow 16 states of charge on the FG

We then can use different voltages on the CG to differentiate these charges

This is really an ***analog*** way of storing bits in memory

Density goes up

But so does ***block*** size, error rate, wear out and access time

Why?

NAND Flash Memory P/E Cycles

Here are some P/E cycles for NAND flash before it becomes unreliable:

SLC NAND: $\sim 10^5$ P/E cycles

MLC NAND: $\sim 10^4$ P/E cycles

TLC NAND: $\sim 10^3$ P/E cycles

QLC NAND: ~ 500 P/E cycles

P/E cycles

Recently these P/E cycles have been experimentally increased 10 to 50 times by local thermal annealing during block erasure

*Reading and writing 0s are done a **page** at a time*

Cannot write 1s to a page

*Erasing (writing 1s) must be done a **block** at a time*

1s are written to an entire block

More on **pages** and **blocks** in a few slides

NAND Flash Latency

As a function of bits/cell

	SLC	MLC	TLC	HDD	RAM
P/E cycles	100k	10k	5k	*	*
Bits per cell	1	2	3	*	*
Seek latency (μs)	*	*	*	9000	*
Read latency (μs)	25	50	100	2000-7000	0.04-0.1
Write latency (μs)	250	900	1500	2000-7000	0.04-0.1
Erase latency (μs)	1500	3000	5000	*	*
Notes	* metric is not applicable for that type of memory				


Note that latencies of NAND is far greater than HDD, but less than RAM

Note also that latencies increase as cell density increases

The SSD controller has to be more discriminating regarding voltages coming out of the NAND. Higher densities requires the memory controller to use more precise voltage measurements, increasing latencies

NAND Density Effects

P/E Cycles & Cost / GB



Read Thy Neighbor

Reading NAND flash memory can cause both the cells being read and also adjacent cells to change over several hundred thousand read cycles

Because erasing requires a high voltage

The problem is mitigated by erasing at the block level rather than the page level

Controllers must keep track of cell read counts and rewrite cells when the count reaches a certain level

*Rewrite means erasing (setting entire block to 1s) and then restoring an entire **block***

Compared to Magnetic Memory

Magnetic memory has none of the wear problems associated with flash memory

Actually they do, but

It's not really "wear"

The equivalent numbers of cycles is much **much** larger

As areal density increases, magnetic dipole stability decreases

Because of wear and P/E cycles, SSD file systems must function differently than disk-based file systems

NAND Flash Memory Organization

Make It Look Like Mass Storage

NAND flash was targeted at replacing mass storage such as CDs, DVDs and magnetic disk drives

Groups of cells are arranged into **pages**

Each **page** contains a number of bytes

e.g., 512, 1024, 2048 or 4096 data bytes

Plus bytes for ECC (error correcting)

So a flash **page** is sort of like a disk sector

Sort of emulates a rotating magnetic disk drive

*But differs because only **blocks** of **pages** can be erased (all bits = 1)*

By comparison, for rotating disks each sector can be changed and bits can change from 0 → 1 or 1 → 0 without any issues

Pages are arranged into **blocks**

*Each **block** usually contains between 32 and 128 **pages***

SSD Controllers

To make an SSD behave like an HDD, a significant number of functions are integrated into the SSD Controller.

Wear leveling

Garbage collection

Trim

SLC cache management

Many of these capabilities have analogs in the HDD world
SSD manufacturers keep their implementations of the functions under wraps

Affects SSD performance

Can give then competitive edge

Wear Leveling

Error and Wear

Error correcting

Claude Shannon (Bell Labs) developed the error correcting algorithms still used today

Flash drives usually correct for 1 or 3 bit errors per page

But some MLCs (multi-level cells) require 5 or more

Wear leveling – typical scheme


Each time a file is changed and then saved, it is saved to a different location

The old location of the file is marked as having been rewritten but available


A memory location is never written a 2nd time until all the virgin or erased memory has been exhausted

Wear leveling

Device without Wear Leveling



Device with Wear Leveling



Dynamic vs. Static Wear Leveling

Dynamic Wear Leveling



Static Wear Leveling



Available area



Used area



Wear leveling

SSD Garbage Collection

Key Characteristics of SSDs

SSDs using NAND technology

Fast reads

Slow writes

Very slow erases

Pages are made up of cells

Blocks are made up of pages

Smallest unit of information for READ and WRITE operations

Page

Smallest unit of information for ERASE operation

Block

SSD Garbage Collection

Key Characteristics of SSDs

To re-write a cell that already contains data, you must first erase the cell.

But the smallest unit for a write is a page

And a page must be erased before it can be written again

And the smallest unit for an erase is a block


So, you must erase an entire block to re-write a NAND cell

Recall that erases are the slowest operation

For fast operations, you need to keep a ready pool of ready-to-write blocks

This is the reason for Garbage Collection

SSD Garbage Collection




Data are written to the 9 pages of Block A.

After the write operation is done, Block A's 9 pages are full.

Pages a-c's data are deleted, but pages cannot be individually erased

They are marked as unreadable, but cannot be written to again.


SSD Garbage Collection



Block A and Block B both have invalid pages
(unreadable)


Designed as dark grey areas

SSD Garbage Collection



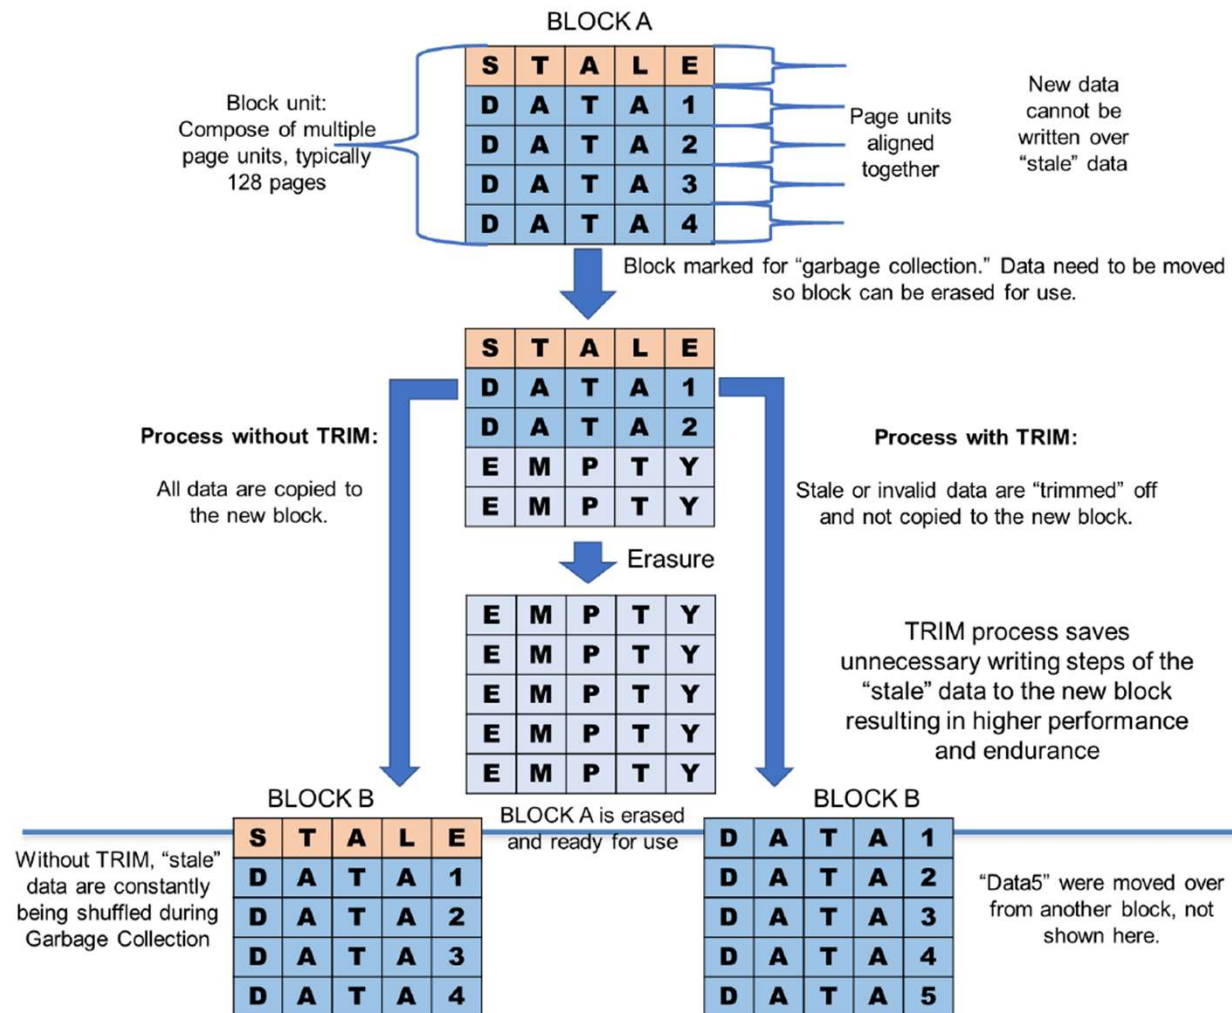
The system will read Pages a-h as containing valid information and transfer them to Block C

SSD Garbage Collection



Garbage collection then erases Blocks A and B
They become available for future writes

SSD Trim



SLC Cache Management

Many SSD Controllers have SLC Caches

These caches are implemented as a portion of the drive

It acts as an SLC with better performance

Two primary types of SLC Caches

Static

- Assigned area on SSD is fixed

- Simpler scheme, but area can wear out more quickly

- Cache size is guaranteed

Dynamic

- Area on SSD not fixed

- Allows for longer lifetime (via wear leveling)

- Cache size isn't guaranteed

SSD Challenges to Forensic Examiner

To summarize

SSDs have sophisticated controllers to emulate HDDs and maximize performance, cost and device lifetime.

This has the following implications for the forensic examiner

On HDDs, when a file is deleted

The space is marked as free

The data will remain there indefinitely until the space is needed for another write operation

On SSDs, when a file is deleted

The space taken by the file is also marked as free

Garbage collection will eventually erase the deleted data

Familiar File Systems

A Quick Overview

Some Common File Systems

FATfs (*File Allocation Table*)

Several variants

FAT12, FAT16, FAT32, VFAT...

Used in rotating mass storage

Also used in SD cards, USB devices

Really flash memory that is made to appear as a FATfs

NTFS (*New Technology File System*)

Developed by Microsoft

Used in all Windows OSs

Several variants with small differences

Some Common File Systems

EXTfs (*Extended File System*)

Several variants (EXT2, EXT3, EXT4)

Used as default in most Linux systems

Flash File Systems

NAND Flash Needs

Because of the special issues surrounding NAND flash memory, file systems tailored to NAND flash are needed

Some issues

*Flash memory **blocks** must be erased before they are written*

To minimize **block** erasures, some schemes cleverly try to write only where existing 1s can be changed to 0s

e.g., 0101 -> 0100 -> 0000

The changes require only changes from 1 -> 0

Rotating disks are not random access. Flash is!

Wear leveling is highly desirable.

Some Well Known Flash File Systems

TrueFFS

The first FFS; 1992

FFS2

Late 1992. Designed for use with MSDOS

Did some wear leveling

FTL

1994

Mapped LBA sector numbers to NAND flash pages

Wear leveling

Multiple proposed enhancements over the years

Some Well Known Flash File Systems

LogFS

Similar to JFFS2, but better at handling large memories in the multi-gigabyte range

We'll discuss JFFS2 next

JFFS2

JFFS2 released in 2001

First Linux FFS

Included in Linux kernels since v. 2.4.10

JFFS2 blocks that are the same size as the NAND flash **blocks**

The smallest number of bits that can be erased (set to 1s)

JFFS2 “nodes” are contained in NAND **blocks**

Nodes contain items such as files and directories

Each node has an inode containing metadata about the node

Nodes consist of one or more pages

Blocks \supset nodes \supset pages

JFFS2

Allocation

As files are updated, a new copy of the file is located in heretofore unused nodes or nodes that are part of cleaned blocks

Journaling

Log of JFFS2 “nodes”

*States of JFFS2 **nodes***

Valid: When, as yet, unused or previously erased (all 1s)

Obsolete: Happens when a newer version of the contents of the node is located somewhere else

JFFS2

Blocks are filled with nodes from low numbered blocks to high numbered blocks

Block States

Clean: All nodes are valid

Dirty: Block contains at least one obsolete node

Free: Has no nodes, valid or obsolete

Garbage Collector

Background job

Erases dirty blocks, converting them into free blocks

If dirty block has both obsolete and valid nodes, it copies the valid nodes to another location before erasing

JFFS2

Wear leveling

Copying updated files to a valid node affects wear leveling

Also, as needed, contents of clean blocks containing fixed but heavily read info are copied to another block

YAFFS2

Yet Another Flash File System

YAFFS has many similarities to JFFS

Accommodates newer NAND flash needs

YAFFS1 didn't need to write only once before erasing

Note: YAFFS 1 & 2 refer to NAND pages as “chunks” and obsolete pages as “*stale*”

YAFFS2 keeps a database in RAM of the physical location of each chunk (page) and their state

Sort of FAT's FAT or Ext's Block Bitmap

Stale chunks (pages) are marked by an indicator in a reserved byte of the chunk (page)

YAFFS2

As files are updated, a new copy of the file is located in heretofore unused chunk (page)

As each new block is used, it is given a sequence number

The block sequence number combined with the location in the block yields the chunk (page) sequence number

So if files have the same metadata, the latest one can be determined. The earlier ones are stale

Deleted files by unlinking it in the RAM database

YAFFS2 Issues

The file system index is stored in RAM

Means that it must be rebuilt upon each boot

This requires scanning the entire NAND flash memory

Takes a lot of time for multi-gigabyte memories

Write-thru caches

Saved files are written to NAND flash

*This means that a number of small changes might
consume a lot of NAND flash memory*

Newer FFSs use write-back caching

Some issues exist respect to recovery from abrupt
power outages

Beyond YAFFS2

The latest FFSs claim to overcome some of the issues with JFFS2 and YAFFS2

ExtremeFFS SanDisk

Better write efficiency and speed for MLC memory

TargetFFS-NAND

Targets embedding in backing stores

SSD Wrap Up

So now I hope that:

You have an idea of the complexity of flash file systems (FFSs)

How FFSs differ from rotating disc drive file systems

How much more is needed in a FFS

And we haven't included

ECC (Error Correction Codes)

Refreshing

Wrap-Up

Because of

Tunneling wear

Asymmetry of program and erase

And also adjacent read wear

Flash mass storage and file systems are quite a bit different than rotating disk magnetic mass storage and their file systems

This lecture attempts to give you an understanding of what flash file systems do and why

Wrap-Up

But today all the P/E cycle, wear leveling, error managing, etc. is handled in an SSD with on-board electronics

It's “invisible” and almost inaccessible to the computer or software

So today's SSD looks like a rotating disk drive to the computer

It follows the ATA command set standard – usually SATA

File systems such as FAT, EXT and NTFS along with MBRs and PBRs are what the BIOS and operating system “sees”.