File Systems and FAT File Systems

FAT: File Allocation Table

Nelson Chapter 5, pp 209-212

Carrier, Chapter 9, pp 211-227

Carrier, Chapter 8 (Carrier's Basic Reference Model) not covered but worthwhile.

Some General Comments

A significant technology contribution that Carrier makes is his ***Basic Reference Model*** for file systems

When studying file systems and using TSK, it helps to understand this model

But some file systems cannot be easily mapped into Carrier's model

Notable among these are FAT file systems

Some General Comments

In Chapters 9 & 10 Carrier tried to “force” FAT file systems into his *Basic Reference Model*

In my opinion doing this makes these chapters confusing

Nelson, unfortunately, gives little attention to FAT

About 4 pages

We will try, in the slides, to make FAT easier to understand

But before we look at FAT file systems, there are a few other topics to discuss

File systems generally

Kilobytes…

Clusters

File Systems

Reprise: Disk Formatting

There are three things that must be done to a disk in the order shown

1st: Physical formatting

Often called "**low level formatting**"

Done on entire disk

2nd: Partitioning

Conceptually separates the disk into disjoint non-overlapping parts

3rd. Logical formatting

Often called "**high level formatting**"

Done on each partition separately

Logical Formatting and File Systems

Logical formatting involves configuring a file system in a partition

A file system consist of an organization and structure for storing & managing data

Main functions of a file system

- *Tracking free space*
- *Allocating free space*
- *Creating and maintaining directories (i.e., folders)*
- *Creating and maintaining directory and file names*
- *Keeping information as to where the parts of each file are physically stored on the disk*

File Systems

Many different file systems

Different file systems usually work with different OSs

File Systems Commonly Used by Windows

FAT12	(early DOS, diskettes)
FAT16	(later DOS, all Windows OSs)
FAT32	(All Windows OSs that I know of)
exFat	(Win Vista and beyond)
VFAT	(FAT32 with long names (Used in Win95, 98, ME)
FATX	(Xbox (But can be read by any current Win OS)
NTFS	(windows NT, 2000, XP, 2003, win7…win11)

File Systems Commonly Used by Unix, Linux & MacOSs

HPFS	(OS/2, Windows NT)
NWFS	(NetWare)
Ext2, Ext3, Ext4	(Linux)
UFS1, UFS2, ZFS	(SunOS)
HFS	(MacOS up to 8)
HFS+	(MacOS 8.1 and later)
Reiser	(Linux, SUSI Linux)
Others as well	

Kilobytes…

How Big?

How big is a Kilobyte?

$10^3 = 1{,}000$ bytes ?

$2^{10} = 1{,}024$ bytes ?

How big is a megabyte?

$10^6 = (10^3)^2 = 1{,}000{,}000$ bytes ?

$2^{20} = (2^{10})^2 = 1024^2 = 1{,}048{,}576$ bytes ?

$10^3 \times 2^{10} = 1000 \times 1024 = 1{,}024{,}000$ bytes ?

What about a gigabyte?

$10^9 = (10^3)^3 = 1{,}000{,}000{,}000$ bytes ?

$2^{30} = (2^{10})^3 = 1024^3 = 1{,}073{,}741{,}824$ bytes ?

$10^6 \times 2^{10} = 10^6 \times 1024 = 1{,}024{,}000{,}000$ bytes ?

Context

The values used depend on context

If were speaking of transmission bit rates, the conventional mathematical values are used

$Kilobyte/sec. = 10^3 = 1000\ bytes/sec.$

$Megabyte/sec. = 10^6 = 1{,}000{,}000\ bytes/sec.$

$Gigabyte/sec. = 10^9 = 1{,}000{,}000{,}000\ bytes/sec.$

If were discussing computer storage, then the power of two numbers are usually used

$Kilobyte = 2^{10} = 1024\ bytes$

$Megabyte = 2^{20} = 1{,}048{,}576\ bytes$

$Gigabyte = 2^{30} = 1{,}073{,}741{,}824\ bytes$

But Not Always

Often the capacity of removable storage is expressed by mixing the two systems

A 1 gigabyte flash drive could be 10^6 kilobytes

10^6 x 1024 = 1,024,000,000 bytes

Example

A Cruzer 1GB flash drive

Marketed as a 1GByte flash drive

1,002,438,144 byes in 1,957,887 sectors

A flash drive used in RADISH

Marketed as a 32 MByte flash drive

29,696,000 bytes in 58,000 sectors

These don't fit either system

How long would it take?

Suppose that you want to transmit a 1GB file over a 1MB link

How long would it take?

10^9 *Bytes* / 10^6 *Bytes per second* = *1000 seconds?*

NO!

2^{30} *bytes* / 10^6 *bytes/sec.* = *1,073,741,824 / 1,000,000*

= *1,073 seconds*

Clusters

Clusters

All data on a FAT, NTFS, or HPFS partition is stored in allocation units called ***clusters***

Cluster are the smallest unit of disk capacity that can be allocated or deallocated by the file system

It is atomic for the file system

Clusters consist of fixed number of contiguous sectors on a disk

For some file systems cluster size is determined by the size of the partition

The larger the partition, the larger the cluster size

FAT is such a file system

Clusters

If a cluster size is 32 sectors (i.e, 16 KB), even a 1 KB file will consume the entire cluster

Using large cluster sizes leaves a lot of unused space

Comments

Much space is wasted

There's ample places for stuff to hide

FAT File Systems

File Allocation Table File Systems

FAT File Systems
Overview

All FAT file systems have three parts

- *A reserved area*
- *The File Allocation Tables area*
- *The data area*
  - Root directory
  - Content area

The term "**FAT**" can refer to either

- *The entire FAT file system (FATfs) or*
- *The File Allocation Tables (FATat) area*
- *This sometimes causes confusion*

FAT File Systems *(FATfs)*
Versions

There are several versions of FATfs, some of which are

- *FAT12*
- *FAT16*
- *FAT32*
- *exFAT (or FATex)*
- *VFAT*
- *FATX*

FATfs are widely used in flash devices of all types

FAT

Where are FAT file systems used?

- *Floppy diskettes (Remember these?)*
- *USB "thumb" drives*
- *Compact "flash" cards*
- *Camera & cell phone SD memory*
- …

Modern OSs use other file systems

- *But all OSs that I know of also read and write FAT file systems*

FAT12

FAT12 is used specifically on floppy disks

Limited amount of file space

Originally designed for MS-DOS 1.0

First Microsoft OS

Used only floppy disks – not on hard disks

Still used for floppies today – if floppies are used

Cluster size is one sector

FAT16

Used on OSs such as

MS-DOS 3.0 and later

Windows95, release 1

NT 3.5 and 4.0

USB "thumb" drives

Partition sizes limited to ≤ 2.02 GB

Cluster size varies with the size of the disk

FAT32

Used with

- *Windows 95, release 2, Windows 98 and ME*
- *Windows 2000 – Win11*
- *USB "thumb" drives*

Usually only be formatted up to 32 GB with Windows

- *Can be formatted beyond this size but cluster sizes are very large*

Number of directory entries ≤ 65,536

exFAT

Becoming widely used for very large flash drives and flash memory; e.g., videos…

Capacity up to 512 TB

Number of directory entries ≤ 270 million

Cluster sizes

7 MB–256 MB	*4 KB*
256 MB–32 GB	*32 KB*
32 GB–256 TB	*128 KB*
> 256 TB	*Not supported*

FAT16, FAT32 & exFAT Cluster Size Comparison

Drive size	Number of sectors	FAT16	FAT32	exFAT
256-511 MB	16	8 KB	4 KB (8 sectors)	4KB up to 256MB 32KB, 257MB-511MB
512 MB-1 GB	32	16 KB	4 KB	32 KB
1-2 GB	64	32 KB	4 KB	32 KB
2-8 GB	8	N/A	4 KB	32 KB
8-16 GB	16	N/A	8 KB (16 sectors)	32 KB
16-32 GB	32	N/A	16 KB	32 KB
More than 32 GB	64	N/A	32 KB	128KB up to 256TB >256TB not supported

Lots of places for stuff to hide

FAT Volume Organization

A FAT file system volume has the following components in order

Carrier's terms

Partition Boot Sector(s) (PBS) ← Reserved Area

File Allocation Table (FAT1)
Mirror of FAT table (FAT2) } FAT Area

Root Folder (or Directory)
Content of other Folders and files } Data Area

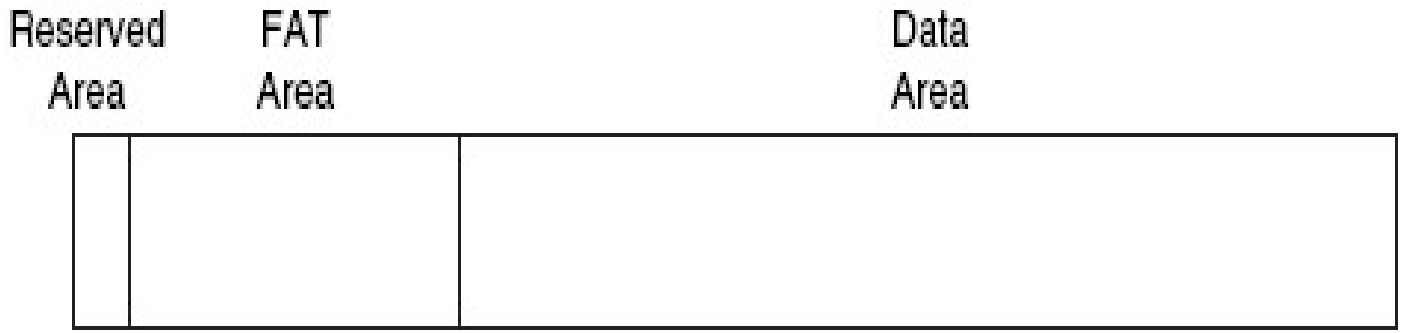

Some Differences Between FAT12/16 and FAT32

FAT32 has somewhat different layout than FAT12 & FAT16

FAT12/16

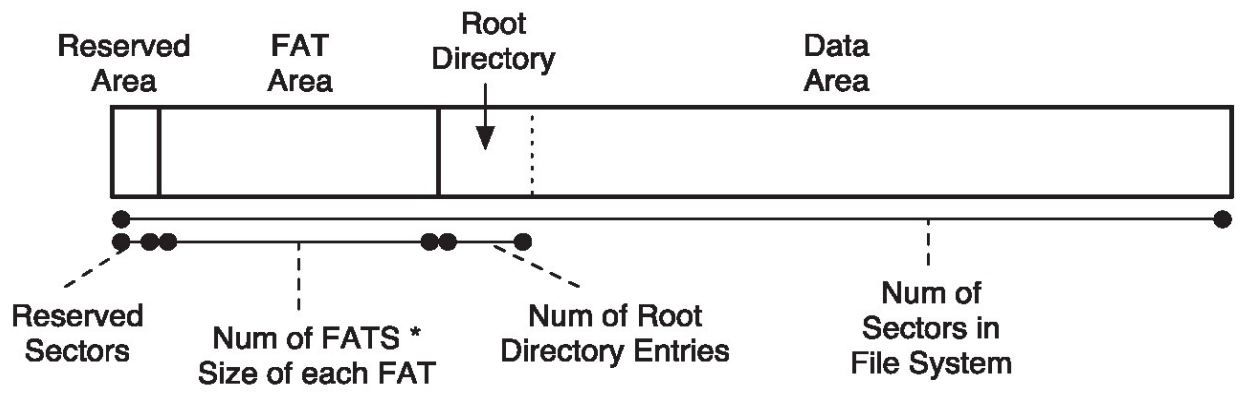

FAT12 &16

One sector in Reserved Area (i.e., PBS)
- *Sometimes > 1 for flash drives*

Root Directory
- Fixed size defined in PBS
- Max of 512 entries
- Located immediately after FAT Area

FAT32

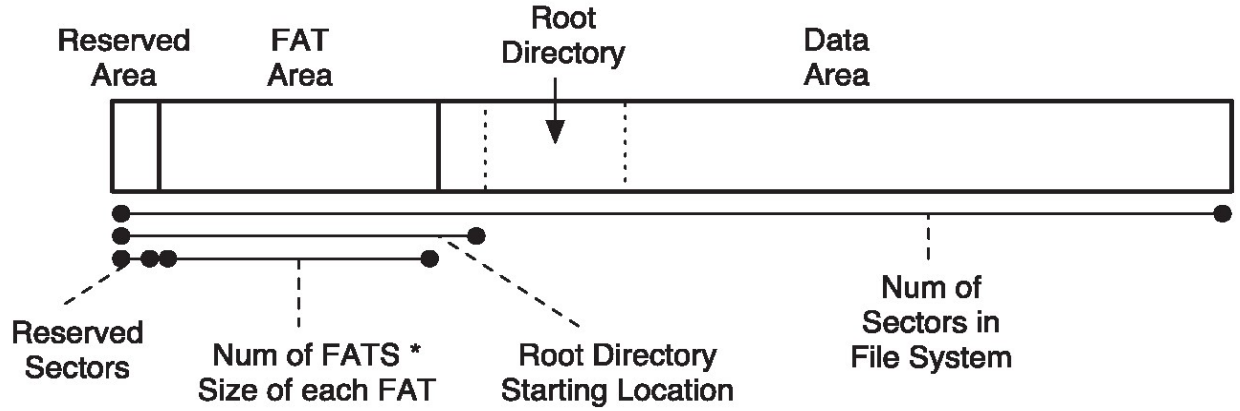

Fat32

Can have multiple sectors in Reserved Area (PBS)

Root Directory
- Size dynamically defined
- Location defined in PBS
- Can be located anywhere in Data Area
- But usually immediately after FAT area

FAT PBS Format

Byte Offset (in hex)	Field Length	Sample Value	Meaning
00	3 bytes	EB 3C 90	Jump instruction
03	8 bytes	MSDOS5.0	OEM Name in text
0B	25 bytes		BIOS Parameter Block
24	26 bytes		Extended BIOS Parameter Block
3E	448 bytes		Bootstrap code
1FE	2 bytes	0x55AA	End of sector marker

FATat
The File Allocation Table Itself

The ***FATfs*** is divided into a fixed number of clusters, each of the same number of sectors

All clusters in a specific **FATfs** *are the same size*

Files in the *data area* (such as user files) are contained in one or more clusters that are not necessarily contiguous

The ***FATat*** (FAT allocation table) is a map of all the clusters in the ***FATfs***

List of entries, one for each cluster in the FATfs

FATat
*The File Allocation Table Itself**

The ***FATat*** contains, for each cluster, one of the following

The cluster number of the next cluster in a chain

An entry indicating the end of cluster chain (EOC)

i.e., the last cluster in a file or folder

An entry to mark a bad cluster

An entry to mark a reserved cluster

An entry to note that the cluster is unused

The number of bits per ***FATat*** entry is

FAT12: *12 bits*

FAT16: *16 bits (2 bytes)*

FAT32: *32 bits (uses 28 bits)*

exFAT: *32 bits*

Contents of FATat Entries
But Not exFAT

FAT12	FAT16	FAT32	Description
0x000	0x0000	0x?0000000	Free Cluster
0x001	0x0001	0x?0000001	Reserved value; do not use
0x002 - 0xFEF	0x0002 - 0xFFEF	0x?0000002 - 0x?FFFFFEF	Used cluster; value points to next cluster
0xFF0 - 0xFF6	0xFFF0 - 0xFFF6	0x?FFFFFF0 - 0x?FFFFFF6	Reserved values; do not use[15].
0xFF7	0xFFF7	0x?FFFFFF7	Bad sector in cluster or reserved cluster
0xFF8 - 0xFFF	0xFFF8 - 0xFFFF	0x?FFFFFF8 - 0x?FFFFFFF	Last cluster in file

The “?” indicates that these 4 bits may be anything.

FAT32 uses only 28 of the 32 bits.

Contents of exFATat Entries*

FAT12	FAT16	FAT32	exFAT	Description
0x000	0x0000	0x?0000000	Not used to indicate free clusters	Free cluster
0x001	0x0001	0x?0000001		Reserved value; do not use
0x002 - 0xFEF	0x0002 - 0xFFEF	0x?0000002 - 0x?FFFFFEF		Used cluster; value points to next cluster
0xFF0 - 0xFF6	0xFFF0 - 0xFFF6	0x?FFFFFF0 - 0x?FFFFFF6		Reserved value; do not use
0xFF7	0xFFF7	0x?FFFFFF7		Bad sector in cluster or reserved cluster
0xFF8 - 0xFFF	0xFFF8 - 0xFFFF	0x?FFFFFF8 - 0x?FFFFFFF		Last cluster in a file

The "?" indicates that these 4 bits may be anything.

FAT32 uses only 28 of the 32 bits.

FAT File/Folder Information Structure

File or folder content is located in clusters pointed to by the “starting cluster number” entry in the FAT

New files are given the first available cluster location on the volume

NT-based OSs can store additional time stamp info

FAT12 & FAT16 Operation

Example: FATat entries for three text files

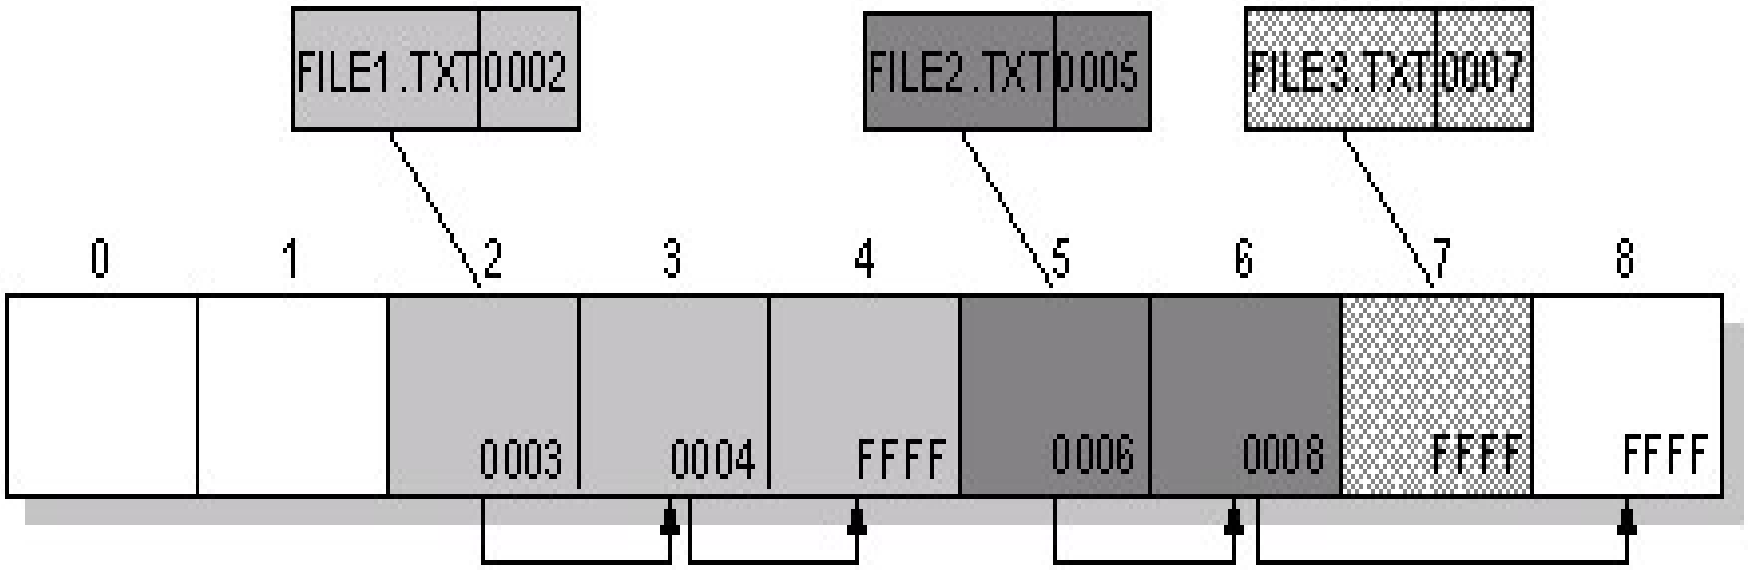

Directory/FAT Operation

Example: FATfs ***Root Directory*** entry

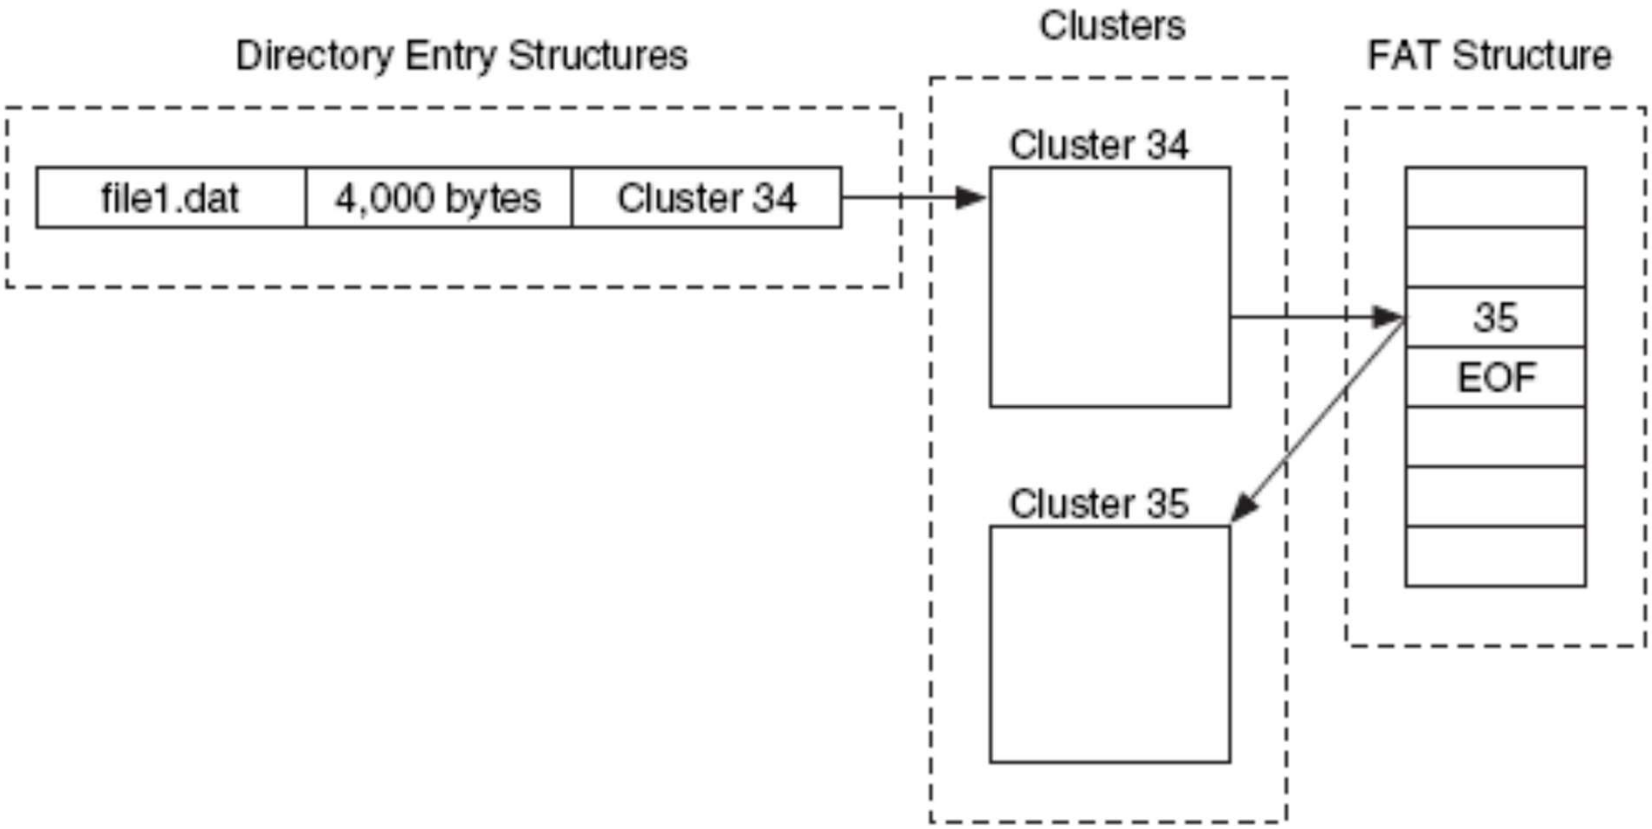

Root Directory

The ***Root Directory*** has a 32-byte entry for each file, folder, subfolder etc. on the FAT partition

File or Folder Name (eight-plus-three characters)

Attribute byte (1 byte)

1 bit each for folder, file, volume label, archive, system, hidden & read-only

Create time (3 bytes)

Create date (2 bytes)

Last access date (2 bytes)

Last modified time (2 bytes)

Last modified date (2 bytes)

Starting cluster number in the file allocation table (2 bytes)

File size (4 bytes) [4.2 GB. But Nelson writes: 2 GB]

Size of File Name

The space for the file or folder name is 8 characters plus a 3 character extension

The file named **forensic.txt** *just fits*

The file name **profnelson.txt** *doesn't quite fit*

And what do we do about .docx

Originally FATfs' had names limited to 8+3 characters

But the file system was modified to allow for much larger names

How?

File Names Longer Than 8+3 Characters

A 2nd *Directory* entry located before the standard *Directory* entry for a file or folder

32 bytes just like normal entry

Can contain up to 26 ASCII or 13 UTF-16 characters

These entries can be chained together to allow for very long names

After the last UTF-16 character a 0x00 0x00 is appended

Not used characters are filled with 0xFF

2nd Directory Entry for Long File Names

This 2nd ***Directory*** entry format

File or Folder name (11 bytes) **Part of long name: 11** bytes

Attribute byte (1 byte)

1 bit each for subfolder, file, volume label, archive, system, hidden, read-only

0x0F (volumeLabel, system, hidden, read-only bits all = 1)

Causes MSDOS to ignore the entry

~~*Create time*~~ *(3 bytes)* **Part of long name: 03** bytes

~~*Create date*~~ *(2 bytes)* **Part of long name: 03** bytes

~~*Last access date*~~ *(2 bytes)* **Part of long name: 02** bytes

~~*Last modified time*~~ *(2 bytes)* **Part of long name: 02** bytes

~~*Last modified date*~~ *(2 bytes)* **Part of long name: 02** bytes

Starting cluster number set to 0x00

~~*File size*~~ *(4 bytes)* **Part of long name: 04** bytes

The bytes used to store a long file name total……: **26** bytes

Analysis of a Root Directory Entry

Beginning of Root Directory

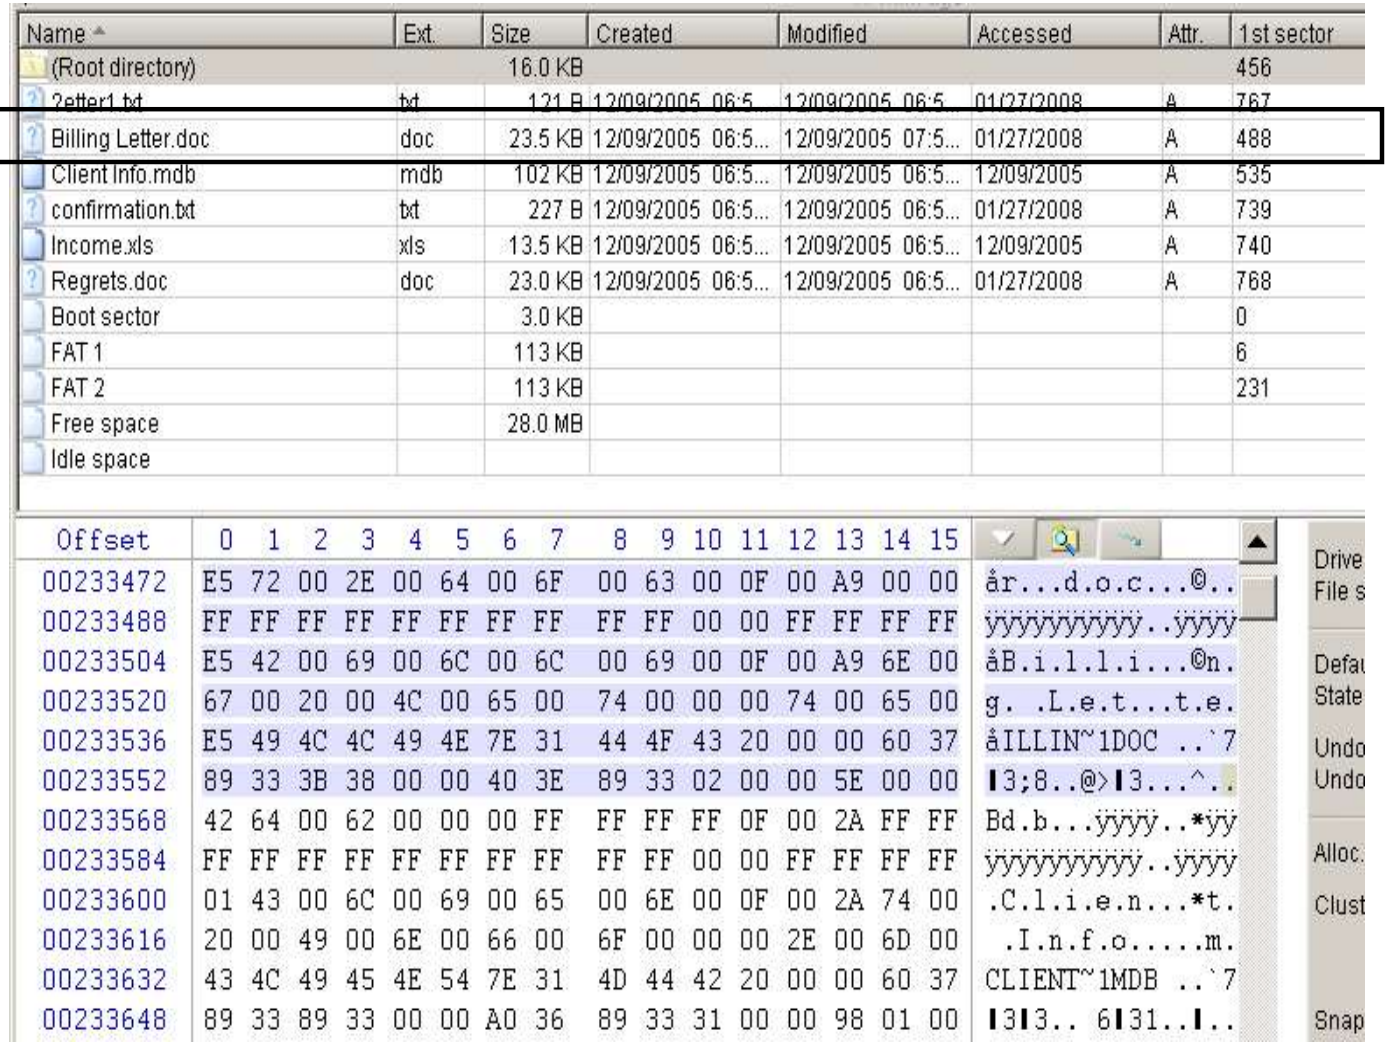

Name	Ext.	Size	Created	Modified	Accessed	Attr.	1st sector
(Root directory)		16.0 KB					456
?etter1.txt	txt	121 B	12/09/2005 06:5...	12/09/2005 06:5...	01/27/2008	A	767
Billing Letter.doc	doc	23.5 KB	12/09/2005 06:5...	12/09/2005 07:5...	01/27/2008	A	488
Client Info.mdb	mdb	102 KB	12/09/2005 06:5...	12/09/2005 06:5...	12/09/2005	A	535
confirmation.txt	txt	227 B	12/09/2005 06:5...	12/09/2005 06:5...	01/27/2008	A	739
Income.xls	xls	13.5 KB	12/09/2005 06:5...	12/09/2005 06:5...	12/09/2005	A	740
Regrets.doc	doc	23.0 KB	12/09/2005 06:5...	12/09/2005 06:5...	01/27/2008	A	768
Boot sector		3.0 KB					0
FAT 1		113 KB					6
FAT 2		113 KB					231
Free space		28.0 MB					
Idle space							

Offset	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
00233472	E5	72	00	2E	00	64	00	6F	00	63	00	0F	00	A9	00	00
00233488	FF	FF	FF	FF	FF	FF	FF	FF	FF	FF	00	00	FF	FF	FF	FF
00233504	E5	42	00	69	00	6C	00	6C	00	69	00	0F	00	A9	6E	00
00233520	67	00	20	00	4C	00	65	00	74	00	00	00	74	00	65	00
00233536	E5	49	4C	4C	49	4E	7E	31	44	4F	43	20	00	00	60	37
00233552	89	33	3B	38	00	00	40	3E	89	33	02	00	00	5E	00	00
00233568	42	64	00	62	00	00	00	FF	FF	FF	FF	0F	00	2A	FF	FF
00233584	FF	FF	FF	FF	FF	FF	FF	FF	FF	FF	00	00	FF	FF	FF	FF
00233600	01	43	00	6C	00	69	00	65	00	6E	00	0F	00	2A	74	00
00233616	20	00	49	00	6E	00	66	00	6F	00	00	00	2E	00	6D	00
00233632	43	4C	49	45	4E	54	7E	31	4D	44	42	20	00	00	60	37
00233648	89	33	89	33	00	00	A0	36	89	33	31	00	00	98	01	00

File Names Longer than 8+3
"Billing Letter.doc"
(14+3)

Name	Ext.	Size	Created	Modified	Accessed	Attr.	1st sector
(Root directory)		16.0 KB					456
?etter1.txt	txt	121 B	12/09/2005 06:5...	12/09/2005 06:5...	01/27/2008	A	767
Billing Letter.doc	doc	23.5 KB	12/09/2005 06:5...	12/09/2005 07:5...	01/27/2008	A	488
Client Info.mdb	mdb	102 KB	12/09/2005 06:5...	12/09/2005 06:5...	12/09/2005	A	535
confirmation.txt	txt	227 B	12/09/2005 06:5...	12/09/2005 06:5...	01/27/2008	A	739
Income.xls	xls	13.5 KB	12/09/2005 06:5...	12/09/2005 06:5...	12/09/2005	A	740
Regrets.doc	doc	23.0 KB	12/09/2005 06:5...	12/09/2005 06:5...	01/27/2008	A	768
Boot sector		3.0 KB					0
FAT 1		113 KB					6
FAT 2		113 KB					231
Free space		28.0 MB					
Idle space							

Offset	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
00233472	E5	72	00	2E	00	64	00	6F	00	63	00	0F	00	A9	00	00
00233488	FF	FF	FF	FF	FF	FF	FF	FF	FF	FF	00	00	FF	FF	FF	FF
00233504	E5	42	00	69	00	6C	00	6C	00	69	00	0F	00	A9	6E	00
00233520	67	00	20	00	4C	00	65	00	74	00	00	00	74	00	65	00
00233536	E5	49	4C	4C	49	4E	7E	31	44	4F	43	20	00	00	60	37
00233552	89	33	3B	38	00	00	40	3E	89	33	02	00	00	5E	00	00
00233568	42	64	00	62	00	00	00	FF	FF	FF	FF	0F	00	2A	FF	FF
00233584	FF	FF	FF	FF	FF	FF	FF	FF	FF	FF	00	00	FF	FF	FF	FF
00233600	01	43	00	6C	00	69	00	65	00	6E	00	0F	00	2A	74	00
00233616	20	00	49	00	6E	00	66	00	6F	00	00	00	2E	00	6D	00
00233632	43	4C	49	45	4E	54	7E	31	4D	44	42	20	00	00	60	37
00233648	89	33	89	33	00	00	A0	36	89	33	31	00	00	98	01	00

- DOS file name
- Attribute (archive)
- Reserved
- Create time (10ms units)
- Create time*
- Create date*
- Last access date
- Used by OS/2
- Last modified time*
- Last modified date*
- First cluster
- File size (bytes)

* Bit encoded

File Names Longer than 8+3
"Billing Letter.doc"
(14+3)

Name	Ext.	Size	Created	Modified	Accessed	Attr.	1st sector
(Root directory)		16.0 KB					456
?etter1.txt	txt	121 B	12/09/2005 06:5...	12/09/2005 06:5...	01/27/2008	A	767
Billing Letter.doc	doc	23.5 KB	12/09/2005 06:5...	12/09/2005 07:5...	01/27/2008	A	488
Client Info.mdb	mdb	102 KB	12/09/2005 06:5...	12/09/2005 06:5...	12/09/2005	A	535
confirmation.txt	txt	227 B	12/09/2005 06:5...	12/09/2005 06:5...	01/27/2008	A	739
Income.xls	xls	13.5 KB	12/09/2005 06:5...	12/09/2005 06:5...	12/09/2005	A	740
Regrets.doc	doc	23.0 KB	12/09/2005 06:5...	12/09/2005 06:5...	01/27/2008	A	768
Boot sector		3.0 KB					0
FAT 1		113 KB					6
FAT 2		113 KB					231
Free space		28.0 MB					
Idle space							

Offset	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
00233472	E5	72	00	2E	00	64	00	6F	00	63	00	0F	00	A9	00	00
00233488	FF	FF	FF	FF	FF	FF	FF	FF	FF	FF	00	00	FF	FF	FF	FF
00233504	E5	42	00	69	00	6C	00	6C	00	69	00	0F	00	A9	6E	00
00233520	67	00	20	00	4C	00	65	00	74	00	00	00	74	00	65	00
00233536	E5	49	4C	4C	49	4E	7E	31	44	4F	43	20	00	00	60	37
00233552	89	33	3B	38	00	00	40	3E	89	33	02	00	00	5E	00	00
00233568	42	64	00	62	00	00	00	FF	FF	FF	FF	0F	00	2A	FF	FF
00233584	FF	FF	FF	FF	FF	FF	FF	FF	FF	FF	00	00	FF	FF	FF	FF
00233600	01	43	00	6C	00	69	00	65	00	6E	00	0F	00	2A	74	00
00233616	20	00	49	00	6E	00	66	00	6F	00	00	00	2E	00	6D	00
00233632	43	4C	49	45	4E	54	7E	31	4D	44	42	20	00	00	60	37
00233648	89	33	89	33	00	00	A0	36	89	33	31	00	00	98	01	00

- DOS file name
- Attribute (archive)
- Reserved
- Create time (10ms units)
- Create time*
- Create date*
- Last access date
- Used by OS/2
- Last modified time*
- Last modified date*
- First cluster
- File size (bytes)
- Long file name
- Reserved
- DOS file name checksum

0x00 indicates end of long file name
Unused characters filled with 0xFF

* Bit encoded

Slack and Chaining

Slack

Slack is the unused space in a cluster not used by the file

There are 2 types of slack

RAM slack

File slack

RAM Slack

The unused space on the last sector of a file

File Slack (sometimes called Drive Slack)

The unused sectors in the cluster in which the file resides

Slack Example
5000 Byte File on a 40 GB FAT Disk

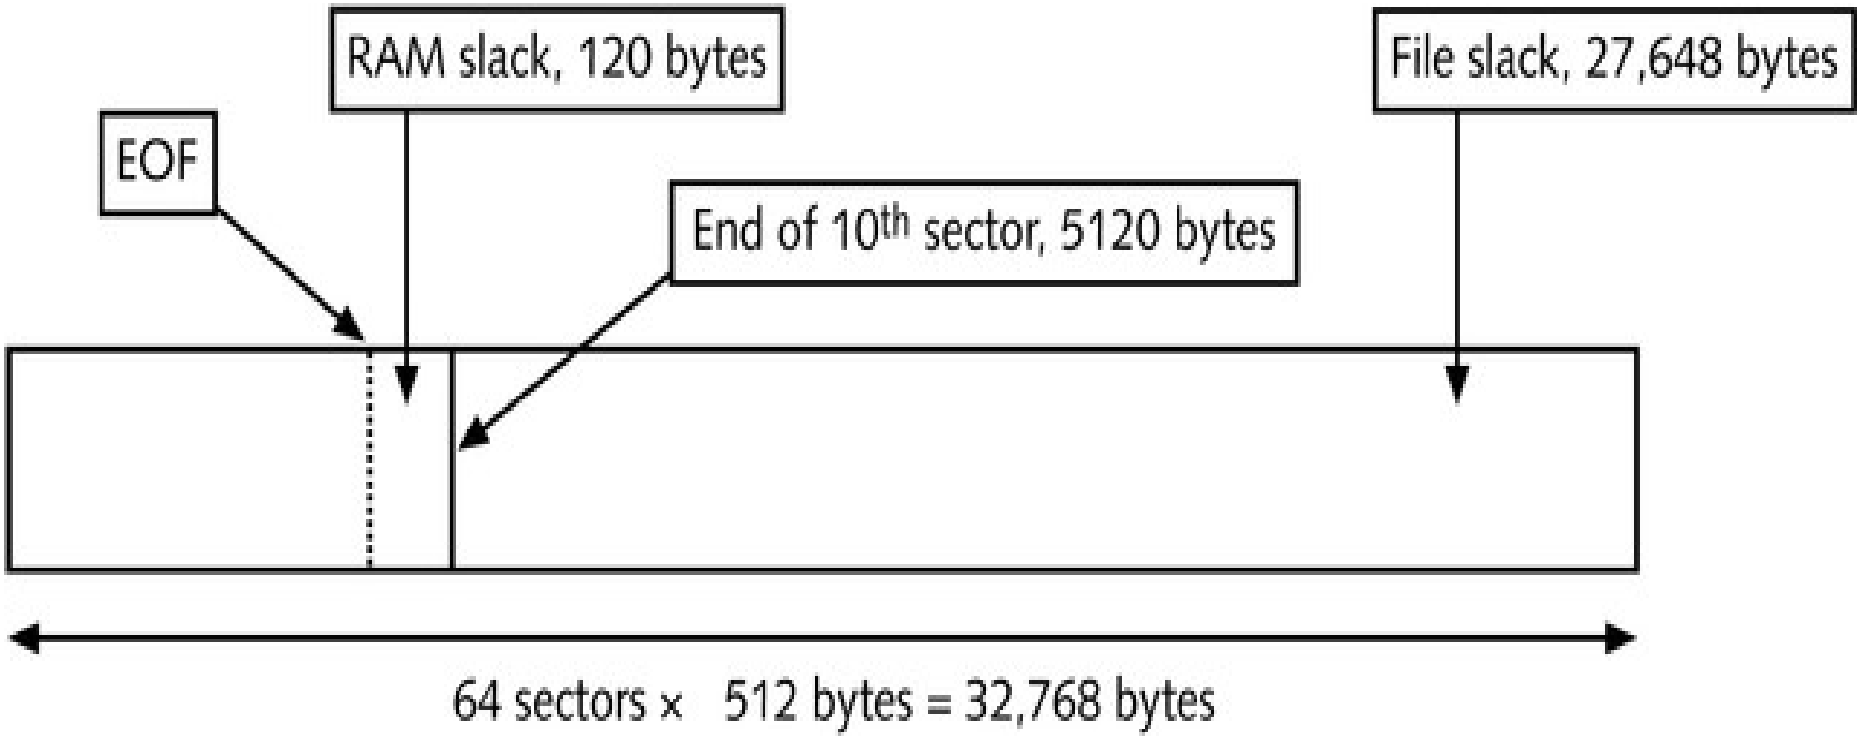

A 5000 byte file occupies 5000/512 or 9.765 sectors

Files Larger Than One Cluster

If a file exceeds the size of a cluster, a 2nd cluster is "chained" to the first cluster

Often these chained clusters are adjacent to each other

But at other times the clusters containing a file are not contiguous

*This is called **fragmentation***

When you defrag a disk, you are rearranging the contents of the disk so that there is minimum fragmentation

Cluster Chaining

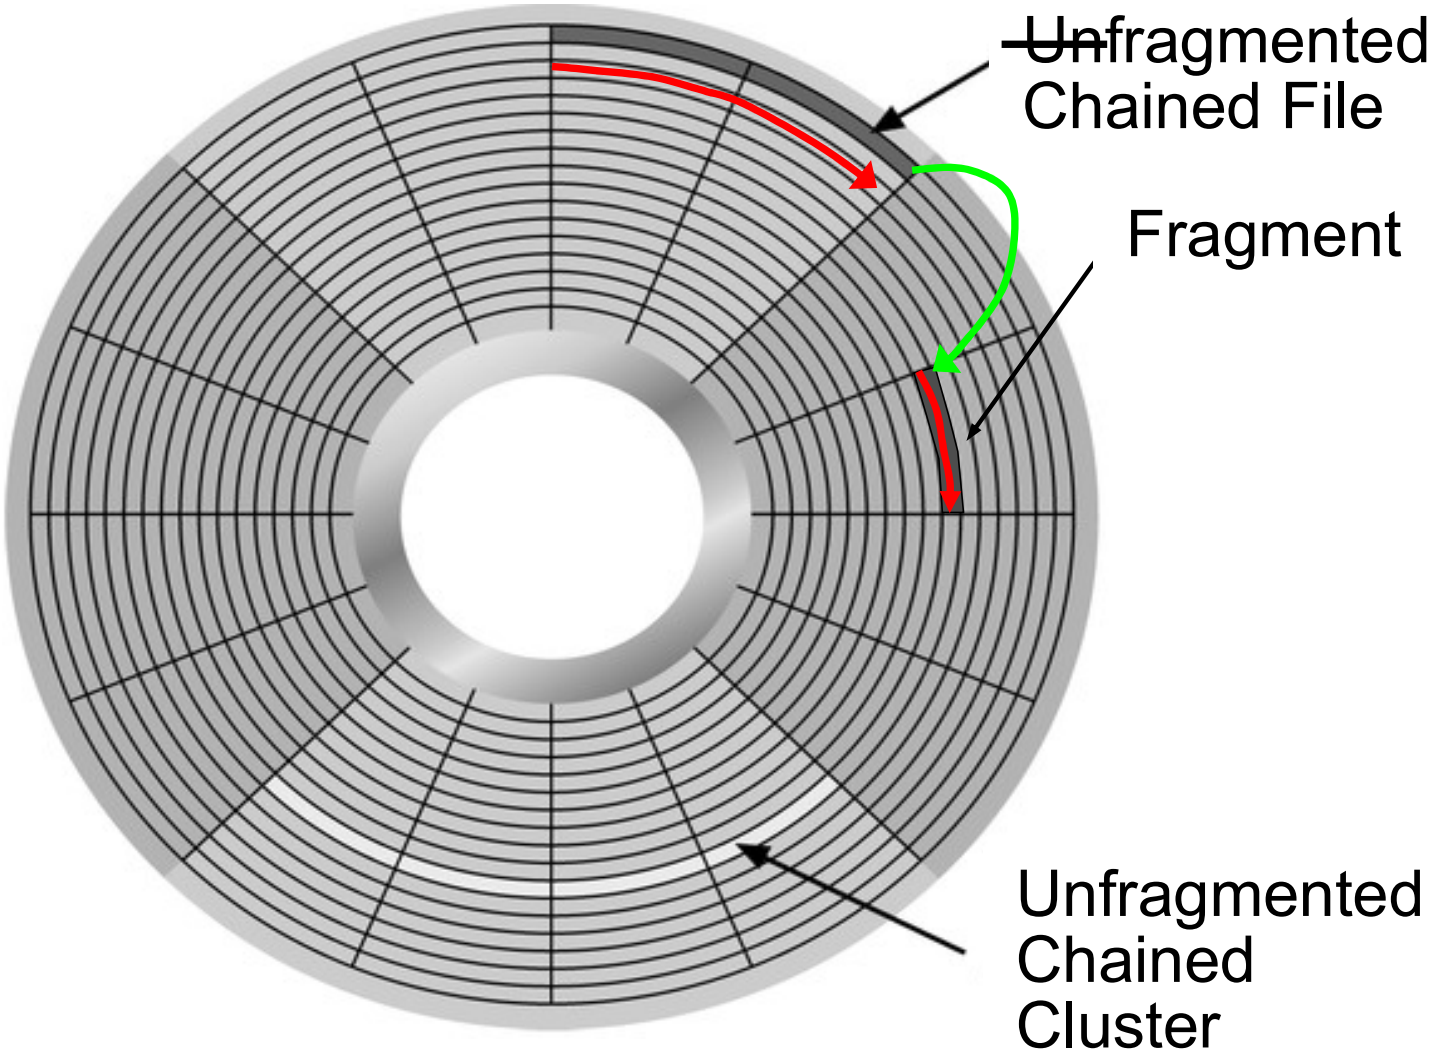

Deleted Files

Deleted Files

In both FAT and NTFS, when a file is deleted

- *The OS replaces the 1st character in the file name with the character* **0xE5**
- *Displayed by many forensic tools as either* **?** *or* **σ**
- *This tells the OS that the file is no longer available*

But it leaves the file intact

- *Forensic and other software tools can recover the file*

The space that the file occupies becomes unallocated

- *Being unallocated, it can be overwritten by other files*
- *In a FATfs the cluster chain is set to zero*